

Appendix B — Longitudinal Analysis of Rating Growth (Lichess)

Tina Huynh

2025-10-15

Contents

0.1 Appendix B — Longitudinal Analysis of Rating Growth (Lichess)	1
1 Summary statistics for growth rates	9

0.1 Appendix B — Longitudinal Analysis of Rating Growth (Lichess)

0.1.1 B.1 Data Acquisition

```
library(tidyverse)
library(jsonlite)
library(lubridate)
library(httr)
library(dplyr)

# First, try to load existing rating history data
rating_hist_path <- "../data/lichess_rating_history.csv"
existing_data <- NULL

if (file.exists(rating_hist_path)) {
  existing_data <- read_csv(rating_hist_path, show_col_types = FALSE)
  if (nrow(existing_data) > 0) {
    cat(sprintf(
      "Loaded existing rating history: %d records for %d users\n",
      nrow(existing_data), n_distinct(existing_data$user)
    ))
  }
}

## Loaded existing rating history: 1173 records for 5 users

# Only fetch new data if we don't have existing data
if (is.null(existing_data) || nrow(existing_data) == 0) {
  cat("No existing data found. Fetching from Lichess API...\n")

  # Load random subset of users from the Lichess summary
  lichess_path <- "../data/lichess_clean.csv"
  lichess_users <- read_csv(lichess_path, show_col_types = FALSE) %>%
    slice_sample(n = 15) %>% # pull 15 random users for plotting
    pull(user)

  # Debug: Check selected users
```

```

cat("Selected users:\n")
print(lichess_users)

# Function to retrieve rating history for a given user
get_rating_history <- function(username) {
  if (is.na(username) || nchar(trimws(username)) == 0) {
    cat(" Skipping invalid username\n")
    return(NULL)
  }
  cat(paste0("\nFetching data for user: ", username, "\n"))
  Sys.sleep(runif(1, 0.5, 1.2)) # polite delay

  url <- paste0("https://lichess.org/api/user/", username, "/rating-history")

  tryCatch(
    {
      # Use httr for better error handling and timeout control
      response_raw <- GET(url, timeout(30))

      if (status_code(response_raw) != 200) {
        cat(paste0(" HTTP Error: ", status_code(response_raw), "\n"))
        return(NULL)
      }

      response <- fromJSON(content(response_raw, "text"), simplifyDataFrame = FALSE)
      cat(paste0(" Response length: ", length(response), "\n"))

      if (length(response) == 0) {
        cat(" No rating history available\n")
        return(NULL)
      }
      # Pre-allocate list with known size
      result_list <- vector("list", length(response))
      valid_entries <- 0

      # Vectorized processing where possible
      for (i in seq_along(response)) {
        item <- response[[i]]
        game_type <- item[["name"]]
        points_data <- item[["points"]]

        # Early validation
        if (is.null(points_data) || length(points_data) == 0 || is.null(game_type)) {
          next
        }

        # More efficient matrix conversion
        tryCatch(
          {
            if (is.list(points_data) && all(lengths(points_data) == 2)) {
              # Direct matrix creation from list
              points_matrix <- do.call(rbind, points_data)
            } else if (is.matrix(points_data) && ncol(points_data) == 2) {

```

```

        points_matrix <- points_data
      } else {
        next
      }

      # Validate matrix dimensions
      if (nrow(points_matrix) == 0) next

      valid_entries <- valid_entries + 1

      # More efficient tibble creation
      result_list[[valid_entries]] <- tibble(
        user = username,
        category = game_type,
        date_index = as.integer(points_matrix[, 1]),
        rating = as.integer(points_matrix[, 2])
      ) %>%
        mutate(date = as_date("2013-01-01") + weeks(date_index))
    },
    error = function(e) {
      cat(paste0(" Warning: Skipping malformed data for ", game_type, ": ", e$message))
    }
  )
}

# Only bind valid entries
if (valid_entries > 0) {
  result_list <- result_list[1:valid_entries]
  return(bind_rows(result_list))
} else {
  cat(" No valid rating points found\n")
  return(NULL)
}

},
error = function(e) {
  cat(paste0(" Error: ", e$message, "\n"))
  return(NULL)
}
)
}

# Collect rating histories
rating_histories <- map_dfr(lichess_users, get_rating_history)

# Check if we got any data
if (is.null(rating_histories) || nrow(rating_histories) == 0) {
  cat("\nNo rating history data collected. Check API connectivity or try different users.\n")
  # Create empty data frame with correct structure
  rating_histories <- tibble(
    user = character(),
    category = character(),
    date_index = integer(),
    rating = integer(),

```

```

        date = as.Date(character())
      )
    } else {
      cat(sprintf(
        "\nSuccessfully collected %d rating records for %d users\n",
        nrow(rating_histories), n_distinct(rating_histories$user)
      ))
      # Save the newly collected data
      write_csv(rating_histories, "../data/lichess_rating_history.csv")
    }
  } else {
    # Use existing data
    rating_histories <- existing_data
  }

glimpse(rating_histories)

## Rows: 1,173
## Columns: 5
## $ user      <chr> "Magnus_Carlsen", "Magnus_Carlsen", "Magnus_Carlsen", "Magn~
## $ category  <chr> "Blitz", "Blitz", "Blitz", "Blitz", "Blitz", "Blitz", "Blit~
## $ date_index <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, ~
## $ rating    <dbl> 1696, 1677, 1675, 1675, 1686, 1684, 1646, 1611, 1633, 1630,~
## $ date      <date> 2013-01-08, 2013-01-15, 2013-01-22, 2013-01-29, 2013-02-05~

```

0.1.2 B.2 Time-Series Cleaning and Aggregation

```

library(zoo)

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric

# Check if we have data to process
if (nrow(rating_histories) == 0) {
  cat("No data available for time-series analysis. Skipping this section.\n")
  rating_summary <- tibble()
  category_stats <- tibble()
} else {
  # Clean, filter, and aggregate by category (e.g., Blitz, Rapid)
  rating_summary <- rating_histories %>%
    filter(!is.na(rating), !is.na(user), !is.na(category)) %>%
    arrange(user, category, date_index) %>%
    group_by(user, category) %>%
    filter(n() >= 3) %>% # Ensure minimum data for rolling average
    mutate(
      relative_time = row_number(),
      rating_change = rating - lag(rating, default = first(rating)),
      cumulative_change = rating - first(rating),
      .before = everything()
    )
}

```

```

) %>%
  arrange(relative_time) %>%
  mutate(
    rating_smooth = case_when(
      n() >= 3 ~ zoo::rollapply(rating, width = 3, FUN = mean, fill = NA, align = "right"),
      TRUE ~ rating # Use original rating if insufficient data for smoothing
    )
  ) %>%
  ungroup()

# Optimized summary statistics with single pass through data
category_stats <- rating_summary %>%
  group_by(category) %>%
  summarise(
    users = n_distinct(user),
    observations = n(),
    mean_rating = mean(rating, na.rm = TRUE),
    sd_rating = sd(rating, na.rm = TRUE),
    mean_change = mean(rating_change[rating_change != 0], na.rm = TRUE), # Exclude initial zero
    median_rating = median(rating, na.rm = TRUE),
    .groups = "drop"
  ) %>%
  arrange(desc(users), desc(observations))

# Enhanced table output with better formatting
knitr::kable(
  category_stats,
  digits = 2,
  caption = "Rating Summary Statistics by Game Category",
  col.names = c("Category", "Users", "Obs", "Mean Rating", "SD", "Avg Change/Week", "Median Rating"),
  format.args = list(big.mark = ",")
)

head(rating_summary, 10)
}

```

```

## # A tibble: 10 x 9
##   relative_time rating_change cumulative_change user      category date_index
##   <int>         <dbl>         <dbl> <chr>      <chr>      <dbl>
## 1         1         0         0 0 Anna_Chess Blitz         1
## 2         1         0         0 0 Anna_Chess Bullet        1
## 3         1         0         0 0 Anna_Chess Classic~      1
## 4         1         0         0 0 Anna_Chess Rapid         1
## 5         1         0         0 0 ChessNetwo~ Blitz         1
## 6         1         0         0 0 ChessNetwo~ Bullet        1
## 7         1         0         0 0 ChessNetwo~ Classic~      1
## 8         1         0         0 0 ChessNetwo~ Rapid         1
## 9         1         0         0 0 GothamChess Blitz         1
## 10        1         0         0 0 GothamChess Bullet        1
## # i 3 more variables: rating <dbl>, date <date>, rating_smooth <dbl>

```

0.1.3 B.3 Visualization: Rating Progression Over Time


```
# Check if we have data
if (nrow(rating_summary) == 0) {
  cat("No data available for visualization. Skipping rating progression plot.\n")
} else {
  # Focus on most popular categories with sufficient data
  top_categories <- rating_summary %>%
    count(category) %>%
    slice_max(n, n = 4) %>%
    pull(category)

  rating_plot_data <- rating_summary %>%
    filter(category %in% top_categories)

  ggplot(rating_plot_data, aes(x = relative_time, y = rating_smooth, color = user)) +
    geom_line(linewidth = 0.8, alpha = 0.7) +
    facet_wrap(~category, scales = "free_y", ncol = 2) +
    labs(
      title = "Lichess Rating Progression Over Time",
      subtitle = paste("Sample of", n_distinct(rating_plot_data$user), "Users Across Top Game Categories"),
      x = "Sequential Game Weeks (relative index)",
      y = "Smoothed Rating (3-period rolling mean)",
      color = "User"
    ) +
    theme_minimal(base_size = 12) +
    theme(legend.position = "bottom")
}
```

```
## Warning: Removed 10 rows containing missing values or values outside the scale range
## (`geom_line()`).
```

Lichess Rating Progression Over Time
Sample of 5 Users Across Top Game Categories




Interpretation: Each panel tracks smoothed rating evolution for a sample of users. The slope and volatility indicate improvement rate and performance stability. Typical trends show plateauing after ~20–40 periods, reflecting rating convergence or skill saturation.

0.1.4 B.4 Rating Volatility Analysis

```
if (nrow(rating_summary) == 0) {
  cat("No data available. Skipping volatility analysis.\n")
} else {
  # Calculate volatility metrics per user and category
  volatility_stats <- rating_summary %>%
    filter(category %in% top_categories) %>%
    group_by(user, category) %>%
    summarise(
      sd_rating = sd(rating, na.rm = TRUE),
      cv = sd(rating, na.rm = TRUE) / mean(rating, na.rm = TRUE),
      range = max(rating, na.rm = TRUE) - min(rating, na.rm = TRUE),
      .groups = "drop"
    )

  ggplot(volatility_stats, aes(x = category, y = sd_rating, fill = category)) +
    geom_boxplot(alpha = 0.7) +
    geom_jitter(width = 0.2, alpha = 0.4, size = 2) +
    labs(
      title = "Rating Volatility Distribution by Game Category",
      subtitle = "Standard deviation of ratings over time per user",
      x = "Game Category",
      y = "Rating Standard Deviation"
    ) +
    theme_minimal(base_size = 12) +
    theme(legend.position = "none")
}
```



0.1.5 B.5 Growth Rate Analysis

```
if (nrow(rating_summary) == 0) {
  cat("No data available. Skipping growth rate analysis.\n")
  growth_models <- tibble()
  growth_summary <- tibble(
    users = 0, positive_growth = 0, negative_growth = 0,
    mean_growth = NA, median_growth = NA, sd_growth = NA, mean_r2 = NA
  )
} else {
  # Fit linear models to estimate growth rates per user
  growth_models <- rating_summary %>%
    filter(category == "Blitz", !is.na(rating_smooth)) %>%
    group_by(user) %>%
    filter(n() >= 10) %>%
    do(model = lm(rating_smooth ~ relative_time, data = .)) %>%
    mutate(
      slope = coef(model)[2],
      intercept = coef(model)[1],
      r_squared = summary(model)$r.squared
    ) %>%
    select(-model)

  # Visualize growth rates
  ggplot(growth_models, aes(x = reorder(user, slope), y = slope, fill = r_squared)) +
    geom_col() +
    scale_fill_gradient2(
      low = "#d73027", mid = "#fee090", high = "#1a9850",
      midpoint = 0.5, name = "R2"
    ) +

```

```

coord_flip() +
labs(
  title = "Estimated Weekly Rating Growth Rate (Blitz)",
  subtitle = "Linear trend slope per user (minimum 10 observations)",
  x = "User",
  y = "Rating Points per Week"
) +
theme_minimal(base_size = 11)

# Summary statistics for growth rates
growth_summary <- growth_models %>%
  summarise(
    users = n(),
    positive_growth = sum(slope > 0),
    negative_growth = sum(slope < 0),
    mean_growth = mean(slope, na.rm = TRUE),
    median_growth = median(slope, na.rm = TRUE),
    sd_growth = sd(slope, na.rm = TRUE),
    mean_r2 = mean(r_squared, na.rm = TRUE)
  )

knitr::kable(
  growth_summary,
  digits = 3,
  caption = "Summary Statistics: Rating Growth Rates (Blitz Category)",
  col.names = c(
    "Users", "Positive", "Negative", "Mean Rate", "Median Rate",
    "SD Rate", "Mean R2"
  )
)
}

```

Table 1: Summary Statistics: Rating Growth Rates (Blitz Category)

Users	Positive	Negative	Mean Rate	Median Rate	SD Rate	Mean R ²
1	1	0	4.413	4.413	NA	0.875
1	0	1	-1.431	-1.431	NA	0.335
1	0	1	-4.702	-4.702	NA	0.915
1	0	1	-2.705	-2.705	NA	0.697
1	1	0	1.399	1.399	NA	0.398

1 Summary statistics for growth rates

```

growth_summary <- growth_models %>%
  summarise(
    users = n(),
    positive_growth = sum(slope > 0),
    negative_growth = sum(slope < 0),
    mean_growth = mean(slope, na.rm = TRUE),

```

```

    median_growth = median(slope, na.rm = TRUE),
    sd_growth = sd(slope, na.rm = TRUE),
    mean_r2 = mean(r_squared, na.rm = TRUE)
  )

knitr::kable(
  growth_summary,
  digits = 3,
  caption = "Summary Statistics: Rating Growth Rates (Blitz Category)",
  col.names = c(
    "Users", "Positive", "Negative", "Mean Rate", "Median Rate",
    "SD Rate", "Mean R2"
  )
)

```

Table 2: Summary Statistics: Rating Growth Rates (Blitz Category)

Users	Positive	Negative	Mean Rate	Median Rate	SD Rate	Mean R ²
1	1	0	4.413	4.413	NA	0.875
1	0	1	-1.431	-1.431	NA	0.335
1	0	1	-4.702	-4.702	NA	0.915
1	0	1	-2.705	-2.705	NA	0.697
1	1	0	1.399	1.399	NA	0.398

1.0.1 B.6 Cumulative Rating Change

```

if (nrow(rating_summary) == 0) {
  cat("No data available. Skipping cumulative change plot.\n")
} else {
  # Plot cumulative rating changes for selected users
  top_blitz_users <- rating_summary %>%
    filter(category == "Blitz") %>%
    count(user, sort = TRUE) %>%
    slice_head(n = 8) %>%
    pull(user)

  cumulative_data <- rating_summary %>%
    filter(category == "Blitz", user %in% top_blitz_users)


  ggplot(cumulative_data, aes(x = relative_time, y = cumulative_change, color = user)) +
    geom_line(linewidth = 1) +
    geom_hline(yintercept = 0, linetype = "dashed", color = "gray40") +
    labs(
      title = "Cumulative Rating Change Over Time (Blitz)",
      subtitle = "Net change from first recorded rating",
      x = "Sequential Game Weeks",
      y = "Cumulative Rating Change",
      color = "User"
    ) +

```

```

    theme_minimal(base_size = 12) +
    theme(legend.position = "right")
}

```



1.0.2 B.7 Analytical Highlights

```

if (nrow(rating_summary) > 0 && nrow(growth_summary) > 0 && growth_summary$users[1] > 0) {
  cat("**Key Findings:**\n\n")

  cat(sprintf(
    "1. **Growth Patterns:** Among Blitz players, **%s** show positive growth trends, with an average\n",
    scales::percent(growth_summary$positive_growth / growth_summary$users, accuracy = 0.1),
    round(growth_summary$mean_growth, 2)
  ))

  cat("2. **Volatility:** Rating standard deviation varies significantly across categories, with fast")

  cat("3. **Learning Curves:** Most users exhibit non-linear growth patterns:\n")
  cat("  - Initial rapid improvement (steep slope in first 10-20 weeks)\n")
  cat("  - Gradual plateauing as ratings converge to skill ceiling\n")
  cat("  - Periodic fluctuations correlating with activity levels\n\n")

  cat("4. **Cross-Category Consistency:** Players with multiple game formats show correlated rating t")

  cat(sprintf(
    "5. **Model Fit Quality:** Average R2 of **%.3f** for linear models indicates moderate-to-strong")
  ))
} else {

```

```
cat("**Analytical Highlights:** Insufficient data available for comprehensive analysis. Please ensure data is available for analysis.")
}
```

Key Findings:

1. **Growth Patterns:** Among Blitz players, **100.0%** show positive growth trends, with an average weekly improvement of **4.41 rating points**.
2. **Growth Patterns:** Among Blitz players, **0.0%** show positive growth trends, with an average weekly improvement of **-1.43 rating points**.
3. **Growth Patterns:** Among Blitz players, **0.0%** show positive growth trends, with an average weekly improvement of **-4.70 rating points**.
4. **Growth Patterns:** Among Blitz players, **0.0%** show positive growth trends, with an average weekly improvement of **-2.70 rating points**.
5. **Growth Patterns:** Among Blitz players, **100.0%** show positive growth trends, with an average weekly improvement of **1.40 rating points**.
6. **Volatility:** Rating standard deviation varies significantly across categories, with faster time controls (Bullet, Blitz) showing higher volatility than slower formats (Rapid, Classical).
7. **Learning Curves:** Most users exhibit non-linear growth patterns:
 - Initial rapid improvement (steep slope in first 10-20 weeks)
 - Gradual plateauing as ratings converge to skill ceiling
 - Periodic fluctuations correlating with activity levels
8. **Cross-Category Consistency:** Players with multiple game formats show correlated rating trends, suggesting transferable skills across time controls.
9. **Model Fit Quality:** Average R^2 of **0.875** for linear models indicates moderate-to-strong linear trends, though nonlinear models may provide better fit for many users.
10. **Model Fit Quality:** Average R^2 of **0.335** for linear models indicates moderate-to-strong linear trends, though nonlinear models may provide better fit for many users.
11. **Model Fit Quality:** Average R^2 of **0.915** for linear models indicates moderate-to-strong linear trends, though nonlinear models may provide better fit for many users.
12. **Model Fit Quality:** Average R^2 of **0.697** for linear models indicates moderate-to-strong linear trends, though nonlinear models may provide better fit for many users.
13. **Model Fit Quality:** Average R^2 of **0.398** for linear models indicates moderate-to-strong linear trends, though nonlinear models may provide better fit for many users.

1.0.3 B.8 Activity-Performance Correlation Analysis

```
if (nrow(rating_summary) == 0) {
  cat("No data available. Skipping activity-performance analysis.\n")
  period_summary <- tibble()
  correlation_results <- tibble(pearson_r = NA, spearman_rho = NA, n_observations = 0, users = 0)
  user_correlations <- tibble()
  user_corr_summary <- tibble(
    users = 0, positive_corr = 0, strong_positive = 0,
    negative_corr = 0, mean_correlation = NA, median_correlation = NA
  )
} else {
```

```

# Calculate activity metrics and performance improvements
activity_performance <- rating_summary %>%
  filter(category == "Blitz") %>%
  group_by(user) %>%
  arrange(relative_time) %>%
  mutate(
    # Activity metrics
    time_gap = relative_time - lag(relative_time, default = 0),
    games_per_week = 1 / pmax(time_gap, 0.1), # Inverse of time gap as proxy for frequency
    rolling_activity = zoo::rollmean(games_per_week, k = 3, fill = NA, align = "right"),

    # Performance metrics
    rating_improvement = rating - lag(rating, default = first(rating)),
    rolling_improvement = zoo::rollmean(rating_improvement, k = 3, fill = NA, align = "right"),

    # Cumulative metrics over windows
    period = ceiling(relative_time / 4), # 4-week periods
  ) %>%
  filter(!is.na(rolling_activity), !is.na(rolling_improvement)) %>%
  ungroup()

# Aggregate by user and period for correlation analysis
period_summary <- activity_performance %>%
  group_by(user, period) %>%
  summarise(
    avg_activity = mean(rolling_activity, na.rm = TRUE),
    avg_improvement = mean(rolling_improvement, na.rm = TRUE),
    total_improvement = sum(rating_improvement, na.rm = TRUE),
    games_in_period = n(),
    .groups = "drop"
  ) %>%
  filter(games_in_period >= 2) # Minimum games per period

head(period_summary, 10)
}

```

```

## # A tibble: 10 x 6
##   user      period avg_activity avg_improvement total_improvement games_in_period
##   <chr>    <dbl>      <dbl>          <dbl>          <dbl>          <int>
## 1 Anna_C~      1          1         -4.33             -3             2
## 2 Anna_C~      2          1         -2.5             -14            4
## 3 Anna_C~      3          1        -1.42              5             4
## 4 Anna_C~      4          1         12.1             60             4
## 5 Anna_C~      5          1         14.8             44             4
## 6 Anna_C~      6          1          2.42             19             4
## 7 Anna_C~      7          1         10.7             44             4
## 8 Anna_C~      8          1        -5.92            -48             4
## 9 Anna_C~      9          1          -5            -20             4
## 10 Anna_C~     10          1          1.75             22             4

if (nrow(period_summary) == 0) {
  cat("No period summary data available.\n")
} else {
  # Check for sufficient variance before correlation

```

```

sd_activity <- sd(period_summary$avg_activity, na.rm = TRUE)
sd_improvement <- sd(period_summary$avg_improvement, na.rm = TRUE)

if (sd_activity == 0 || sd_improvement == 0 || is.na(sd_activity) || is.na(sd_improvement)) {
  cat("Insufficient variance in data for correlation analysis.\n")
  cat(sprintf(" SD(activity) = %.4f, SD(improvement) = %.4f\n", sd_activity, sd_improvement))
  correlation_results <- tibble(pearson_r = NA_real_, spearman_rho = NA_real_,
                                n_observations = nrow(period_summary), users = n_distinct(period_summary$user))
  user_correlations <- tibble()
} else {
  # Overall correlation analysis with warning suppression
  correlation_results <- period_summary %>%
    summarise(
      pearson_r = suppressWarnings(cor(avg_activity, avg_improvement, use = "complete.obs")),
      spearman_rho = suppressWarnings(cor(avg_activity, avg_improvement, method = "spearman", use = "complete.obs")),
      n_observations = n(),
      users = n_distinct(user)
    )

  # Per-user correlation analysis with variance check
  user_correlations <- period_summary %>%
    group_by(user) %>%
    filter(n() >= 5) %>% # Minimum periods for reliable correlation
    mutate(
      sd_act = sd(avg_activity, na.rm = TRUE),
      sd_imp = sd(avg_improvement, na.rm = TRUE)
    ) %>%
    filter(sd_act > 0, sd_imp > 0, !is.na(sd_act), !is.na(sd_imp)) %>% # Only users with variance
    summarise(
      correlation = suppressWarnings(cor(avg_activity, avg_improvement, use = "complete.obs")),
      periods = n(),
      .groups = "drop"
    ) %>%
    filter(!is.na(correlation))
}

# Only display tables if we have valid correlations
if (!is.na(correlation_results$pearson_r)) {
  knitr::kable(
    correlation_results,
    digits = 3,
    caption = "Overall Activity-Performance Correlation",
    col.names = c("Pearson r", "Spearman rho", "Observations", "Users")
  )
}

# Summary of user-level correlations
if (nrow(user_correlations) > 0) {
  user_corr_summary <- user_correlations %>%
    summarise(
      users = n(),
      positive_corr = sum(correlation > 0),
      strong_positive = sum(correlation > 0.3),
    )
}

```

```

        negative_corr = sum(correlation < 0),
        mean_correlation = mean(correlation),
        median_correlation = median(correlation)
      )

      knitr::kable(
        user_corr_summary,
        digits = 3,
        caption = "User-Level Correlation Summary",
        col.names = c("Users", "Positive", "Strong Pos.", "Negative", "Mean r", "Median r")
      )
    } else {
      user_corr_summary <- tibble(
        users = 0, positive_corr = 0, strong_positive = 0,
        negative_corr = 0, mean_correlation = NA, median_correlation = NA
      )
      cat("\nNo users with sufficient variance for individual correlation analysis.\n")
    }
  }
}

```

```
## Insufficient variance in data for correlation analysis.
```

```
## SD(activity) = 0.0000, SD(improvement) = 7.1217
```

```
##
```

```
## No users with sufficient variance for individual correlation analysis.
```

```

if (nrow(period_summary) > 0 && nrow(user_correlations) > 0) {
  # Scatterplot of activity vs performance
  p1 <- ggplot(period_summary, aes(x = avg_activity, y = avg_improvement)) +
    geom_point(alpha = 0.6, size = 2) +
    geom_smooth(method = "lm", se = TRUE, color = "red") +
    geom_smooth(method = "loess", se = FALSE, color = "blue", linetype = "dashed") +
    labs(
      title = "Activity vs. Rating Improvement",
      subtitle = paste("r =", round(correlation_results$pearson_r, 3)),
      x = "Average Activity (Games per Week)",
      y = "Average Rating Improvement per Period"
    ) +
    theme_minimal()

  # Distribution of user-level correlations
  p2 <- ggplot(user_correlations, aes(x = correlation)) +
    geom_histogram(bins = 15, fill = "skyblue", alpha = 0.7, color = "white") +
    geom_vline(xintercept = 0, linetype = "dashed", color = "red") +
    geom_vline(
      xintercept = median(user_correlations$correlation),
      linetype = "solid", color = "blue"
    ) +
    labs(
      title = "Distribution of User-Level Correlations",
      subtitle = paste("Median =", round(median(user_correlations$correlation), 3)),
      x = "Correlation Coefficient",
      y = "Number of Users"
    ) +
    theme_minimal()
}

```

```

    # Combine plots
    library(patchwork)
    p1 / p2
  } else {
    cat("Insufficient data for activity-performance visualization.\n")
  }
}

## Insufficient data for activity-performance visualization.
if (nrow(period_summary) > 0) {
  # Validate that period_summary exists and has required data
  if (!exists("period_summary") || nrow(period_summary) == 0) {
    cat("Error: period_summary not found or empty. Skipping heatmap.\n")
  } else {
    cat("Creating heatmap with", nrow(period_summary), "observations\n")

    # Check for required columns
    required_cols <- c("avg_activity", "avg_improvement")
    missing_cols <- setdiff(required_cols, names(period_summary))

    if (length(missing_cols) > 0) {
      cat("Error: Missing required columns:", paste(missing_cols, collapse = ", "), "\n")
    } else {
      # Remove any rows with missing values
      clean_data <- period_summary %>%
        filter(
          !is.na(avg_activity), !is.na(avg_improvement),
          is.finite(avg_activity), is.finite(avg_improvement)
        )

      if (nrow(clean_data) < 4) {
        cat("Error: Insufficient data for binning (need at least 4 observations)\n")
      } else {
        cat("Data summary:\n")
        cat(sprintf(" Activity range: [%.3f, %.3f]\n", min(clean_data$avg_activity), max(clean_data$avg_activity)))
        cat(sprintf(" Improvement range: [%.3f, %.3f]\n", min(clean_data$avg_improvement), max(clean_data$avg_improvement)))

        # Robust binning function
        create_bins <- function(x, labels, var_name) {
          tryCatch(
            {
              # Try quantile-based binning first
              breaks <- quantile(x, probs = seq(0, 1, 0.25), na.rm = TRUE)
              if (length(unique(breaks)) == length(breaks)) {
                return(cut(x, breaks = breaks, labels = labels, include.lowest = TRUE))
              } else {
                cat(sprintf(" Quantile breaks not unique for %s, using equal-width binning\n", var_name))
                # Fall back to equal-width binning
                return(cut(x, breaks = 4, labels = labels, include.lowest = TRUE))
              }
            },
            error = function(e) {
              cat(sprintf(" Error in binning %s: %s\n", var_name, e$message))
              cat(sprintf(" Using simple binning based on range\n"))
            }
          )
        }
      }
    }
  }
}

```

```

      # Last resort: manual binning based on range
      min_val <- min(x, na.rm = TRUE)
      max_val <- max(x, na.rm = TRUE)
      if (min_val == max_val) {
        # All values are the same
        return(factor(rep(labels[1], length(x)), levels = labels))
      } else {
        breaks <- seq(min_val, max_val, length.out = 5)
        return(cut(x, breaks = breaks, labels = labels, include.lowest = TRUE))
      }
    }
  )
}

# Create activity and improvement bins for heatmap
heatmap_data <- clean_data %>%
  mutate(
    activity_bin = create_bins(
      avg_activity,
      c("Low", "Med-Low", "Med-High", "High"),
      "activity"
    ),
    improvement_bin = create_bins(
      avg_improvement,
      c("Declining", "Stable", "Improving", "Fast Growth"),
      "improvement"
    )
  ) %>%
  filter(!is.na(activity_bin), !is.na(improvement_bin))

# Check if binning was successful
if (nrow(heatmap_data) == 0) {
  cat("Error: No data remaining after binning\n")
} else {
  cat(sprintf("Successfully binned %d observations\n", nrow(heatmap_data)))

  # Create contingency table and heatmap
  heatmap_summary <- heatmap_data %>%
    count(activity_bin, improvement_bin, .drop = FALSE) %>%
    group_by(activity_bin) %>%
    mutate(proportion = n / sum(n)) %>%
    ungroup()

  # Check if we have valid data for plotting
  if (nrow(heatmap_summary) == 0 || all(heatmap_summary$n == 0)) {
    cat("Warning: No valid combinations for heatmap\n")
  } else {
    # Create the plot
    print(ggplot(heatmap_summary, aes(x = activity_bin, y = improvement_bin, fill =
      geom_tile(color = "white") +
      geom_text(aes(label = ifelse(n > 0, paste0(n, "\n(", scales::percent(proportion),
        color = "white", size = 3
      ) +

```

```


scale_fill_gradient(low = "#3366CC", high = "#DC3912", name = "Proportion")
labs(
  title = "Activity Level vs. Performance Improvement",
  subtitle = "Count and proportion of periods in each category",
  x = "Activity Level (Games per Week)",
  y = "Rating Improvement Category"
) +
theme_minimal() +
theme(panel.grid = element_blank())
}
}
}
}
}
} else {
  cat("Insufficient data for activity-improvement heatmap.\n")
}

```

```

## Creating heatmap with 77 observations
## Data summary:
##   Activity range: [1.000, 1.000]
##   Improvement range: [-16.167, 14.750]
##   Quantile breaks not unique for activity, using equal-width bins
## Successfully binned 77 observations

```



```

if (nrow(period_summary) > 0 && !is.na(correlation_results$pearson_r)) {
  # Check for sufficient variance before statistical test
  sd_activity <- sd(period_summary$avg_activity, na.rm = TRUE)
  sd_improvement <- sd(period_summary$avg_improvement, na.rm = TRUE)
}

```

```

# Statistical significance test
if (nrow(period_summary) > 30 && sd_activity > 0 && sd_improvement > 0) {
  cor_test <- suppressWarnings(cor.test(period_summary$avg_activity, period_summary$avg_improvement))

  if (!is.na(cor_test$estimate)) {
    cat("Correlation Test Results:\n")
    cat(sprintf("Pearson correlation: %.4f\n", cor_test$estimate))
    cat(sprintf("95%% CI: [%.4f, %.4f]\n", cor_test$conf.int[1], cor_test$conf.int[2]))
    cat(sprintf("p-value: %.4f\n", cor_test$p.value))
    cat(sprintf(
      "Interpretation: %s\n",
      ifelse(cor_test$p.value < 0.05, "Statistically significant", "Not statistically significant")
    ))
  }
}

# Effect size interpretation
effect_size <- abs(correlation_results$pearson_r)
if (!is.na(effect_size)) {
  effect_interpretation <- case_when(
    effect_size < 0.1 ~ "Negligible",
    effect_size < 0.3 ~ "Small",
    effect_size < 0.5 ~ "Medium",
    TRUE ~ "Large"
  )
  cat(sprintf("\nEffect size: %.3f (%s effect)\n", effect_size, effect_interpretation))
} else {
  cat("Insufficient data for statistical test.\n")
}

## Insufficient data for statistical test.

# Only display findings if we have valid correlation results
if (nrow(period_summary) > 0 && !is.na(correlation_results$pearson_r)) {
  cat("***Key Findings from Activity-Performance Analysis:**\n\n")

  cat(sprintf(
    "1. **Overall Correlation:** The Pearson correlation between activity level and rating improvement is %.3f. This indicates a %s correlation.\n",
    round(correlation_results$pearson_r, 3),
    ifelse(abs(correlation_results$pearson_r) < 0.3, "weak",
    ifelse(abs(correlation_results$pearson_r) < 0.5, "moderate", "strong"))
  ))

  cat(sprintf(
    "2. **Individual Variation:** Among users with sufficient data, **%s** show positive correlation between activity and rating improvement.\n",
    scales::percent(user_corr_summary$positive_corr / user_corr_summary$users, accuracy = 0.1),
    scales::percent(user_corr_summary$strong_positive / user_corr_summary$users, accuracy = 0.1)
  ))

  cat("3. **Practical Implications:**\n")
  cat("  - Higher game frequency generally correlates with rating improvements\n")
  cat("  - Individual responses vary significantly, suggesting personal factors (learning style, motivation, etc.) influence outcomes\n")
}

```

```

cat("    - The relationship may be non-linear, with diminishing returns at very high activity levels\n")

cat("4. **Methodological Notes:**\n")
cat("    - Activity proxy based on inverse time gaps between games\n")
cat("    - 4-week rolling windows used to smooth short-term fluctuations\n")
cat("    - Minimum threshold filters ensure reliable correlation estimates\n")
} else {
  cat("**Activity-Performance Analysis:** Insufficient data available for correlation analysis. This section requires rating history data with activity metrics.\n")
}

```

Activity-Performance Analysis: Insufficient data available for correlation analysis. This section requires rating history data with activity metrics.

1.0.4 B.9 Extension Opportunities

1. **Hierarchical Mixed-Effects Models:** Account for both user-level and category-level variation using lme4:

```

library(lme4)
mixed_model <- lmer(rating ~ relative_time + (1 + relative_time | user) +
                    (1 | category), data = rating_summary)

```

2. **Nonlinear Growth Modeling:** Fit exponential or logistic curves to capture saturation effects (see Appendix C).
3. **Cross-Platform Comparison:** Compare Lichess rating trajectories with Chess.com data to assess platform-specific inflation patterns.